

Grammar-Derived Operation Abstraction for Cross-Project Vulnerability Detection

Anonymous Authors

Abstract

Deep learning detectors of software vulnerabilities generalise poorly across projects, and the field has responded largely by proposing new architectures. We ask a different question: how much of the cross-project gap is attributable to the *representation* rather than the model? We specify ϕ , a total labelling function from the node alphabet of a C grammar onto a small operation taxonomy, together with a projection lattice that makes its three granularities nested by construction. Because the resulting alphabet is fixed by the grammar rather than by any project’s vocabulary, the representation is identical across corpora by construction. We evaluate ϕ twice. As a static analysis, we annotate a stratified sample of call sites and report per-class precision and recall, which exposes failure modes that an accuracy-only evaluation hides — notably that substring-based recognition of wrapped allocators misclassifies writes and logging macros as memory operations. As a representation, we hold graph topology, backbone, depth, readout and training budget fixed and vary only the node features, under project-level GroupKFold with deduplication, mega-projects held train-only, and cluster-bootstrap confidence intervals over held-out projects. Replacing lexical node features with the abstraction improves cross-project ranking by +0.064 AUC (95% CI [+0.028, +0.112], $p = 0.001$) and +0.022 AUPRC (95% CI [+0.000, +0.049], $p = 0.047$) on BigVul, winning 8 of 10 folds. We report negative results with equal weight: learned edge weighting and affinity-weighted federated aggregation do not help, and refining the taxonomy beyond eight types does not either. All code, annotations and the scripts that regenerate every number are released.

Keywords: Program Representation, Static Analysis, Abstract Syntax Trees, Software Vulnerability Detection, Federated Learning, Graph Neural Networks, Differential Privacy

1. Introduction

Automated vulnerability detection has become a proving ground for graph learning: a function is parsed into a graph, a neural network is trained over it, and detection accuracy is reported. Progress is usually claimed architecturally — a different message-passing scheme, an attention variant, a language-model hybrid. Yet replication studies keep finding the same thing: whatever the architecture, performance collapses when the model is evaluated on projects it has never seen [1, 2, 3]. Detectors that appear strong within a project transfer poorly across projects, and successive architectures have not closed that gap.

This paper examines a factor that most of that work holds constant: what a node in the graph actually *carries*. Nearly all graph-based detectors embed the node’s lexical content — an identifier, an API name, a literal. Such a vocabulary is a property of a codebase, not of a language. FFmpeg allocates with `av_malloc`, GLib with `g_malloc`, the Linux kernel with `kmalloc`, OpenSSL with `OPENSSL_malloc`. A model whose features are drawn from one project’s vocabulary has, at test time on another project, largely out-of-vocabulary input. The mechanism the model must learn — an allocation whose size is computed, an indexed write, a missing bound — is shared; the tokens denoting it are not.

We therefore study a program abstraction rather than an architecture. We define ϕ , a total labelling function from the node-kind alphabet of the tree-sitter C grammar onto an operation taxonomy of eight types, with finer 16- and 32-type refinements related by explicit surjections so that the three granularities are nested by construction (Proposition 1). The resulting node alphabet is fixed by the grammar and the taxonomy, so it is identical for every project and every corpus — the property we claim is responsible for transfer.

Two evaluations follow, and we regard the first as being as important as the

second. Abstractions of this kind are usually introduced as preprocessing and never validated: no precision, no recall, no error analysis. We instead evaluate ϕ as a static analysis, annotating a stratified sample of call sites and reporting per-class precision and recall. That exercise is not ceremonial. It found that substring-based recognition of wrapped allocators labels `__put_user` and `skb_put` — both writes — as release operations, treats a floating-point predicate as a block set, and misses `copy_to_user`; and that classifying `printf` and `fprintf` as buffer-writing formatters silently counts 655 stream-I/O call sites in one corpus as memory operations. Both defects are invisible to downstream accuracy and were corrected before the experiments reported here.

The second evaluation isolates the representation. Graph topology, backbone, depth, hidden width, readout, optimiser and gradient budget are held fixed; only the node features vary. Evaluation uses project-level GroupKFold with near-duplicate removal before splitting, mega-projects held train-only so that no “cross-project” fold is secretly a single repository, vocabularies fitted on training projects only, and cluster-bootstrap confidence intervals that resample held-out *projects* rather than functions.

Our contributions are:

- **A specified abstraction.** ϕ as a total, deterministic labelling over a grammar alphabet, with a projection lattice and a nesting proposition, requiring no name resolution, build system or whole-program context (Section 3).
- **Validation as an analysis.** Per-class precision and recall against annotated ground truth, with the measured precision/coverage trade-off between exact allow-listing and component-based recognition of project-wrapped APIs, and the failure modes each admits (Section 4.2).
- **The representation effect, measured under controls.** Replacing lexical features with the abstraction improves cross-project ranking by +0.064 AUC ($p = 0.001$) and +0.022 AUPRC ($p = 0.047$), winning 8 of 10 folds, with every other factor held constant (Section 4.3).

• **Negative results, reported with equal weight.** A learned edge-weighting module and an affinity-weighted federated aggregation scheme — both of which we implemented and initially believed in — do not improve on their ablations; nor does refining the taxonomy beyond eight types (Section 4.5).

Section 2 positions the work, Section 3 specifies the abstraction, Section 4 reports both evaluations, and Section 5 discusses what the result does and does not license.

2. Background and Related Work

2.1. Code Representations for Learned Detectors

Learned vulnerability detection has moved through three representations. Early sequence models treated a function as a token stream: VulDeePecker’s code gadgets [4] and SySeVR’s syntax- and semantics-based candidates [5] sliced programs into data-dependent token sequences for recurrent classifiers. The Code Property Graph (CPG) [6], unifying abstract syntax, control flow and data dependence in one structure, then made graph learning natural, and Devign [7] established the template still in use: parse to a graph, embed node text, propagate, pool, classify. Later work varied the graph or the propagation — inter-procedural slices in DeepWukong [8], feature-attention over program dependence in IVDetect [9], heterogeneous attention in [10]. A third line embeds code with pre-trained language models, either alone as in LineVul [11] or fused with a graph channel [12, 13].

What is largely invariant across this progression is the node feature: some embedding of the node’s lexical content. Our work isolates that choice. We hold the graph, the propagation scheme and the training protocol fixed and vary only what a node carries, which — to our knowledge — prior work does not do as a controlled comparison.

2.2. Normalisation and Abstraction of Code

85 Abstracting away identifiers is long-established practice, but as preprocessing
rather than as an object of study. Clone detection normalises identifiers and
literals so that Type-2 clones become textually identical [6]; sequence-based
detectors routinely map user-defined names to placeholder tokens (`VAR1`, `FUN1`)
before training [4, 5]; the Devign corpus itself ships a normalised variant of every
90 function.

Two properties distinguish what we do. First, existing normalisation is
anonymising: it erases the identifier and replaces it with a fresh symbol, pre-
serving only equality of occurrences. Ours is *semantic*: it maps a call onto the
memory effect the callee is documented to have, so that `av_malloc` and `kmalloc`
95 become the same symbol as `malloc` rather than three distinct placeholders. Sec-
ond, and more importantly for a languages venue, we treat the mapping as an
analysis with a specification and a measured error rate, rather than as an un-
examined preprocessing step. We are not aware of prior work in this area that
reports precision and recall for its own abstraction.

100 The idea of typing operations by their effect is of course not new in program-
ming languages — it is the intuition behind effect systems and behavioural types
— but those require type information and whole-program context that function-
level vulnerability corpora do not provide. Our ϕ is deliberately weaker: syntax-
directed, name-based, and total, so it applies to isolated, non-preprocessed func-
105 tions, at the cost of the precision quantified in Section 4.2.

2.3. The Cross-Project Generalisation Gap

The finding that motivates this paper is now well replicated. Chakraborty
et al. [1] showed that detectors reporting strong within-project results degrade
sharply on realistic, project-disjoint data; DiverseVul [14] demonstrated that
110 simply adding training projects does not close the gap; PrimeVul [3] showed
that much apparent progress reflects duplication and label noise, with F1 falling
dramatically under a rigorous split; and [2] found the same degradation at
repository scale.

Responses have been mostly architectural or adaptational — domain adap-
tation across projects [15, 16, 17], transfer with metric fusion [18]. Our result is
115 complementary and, we argue, prior: with architecture held constant, changing
the node representation alone yields a larger and more consistent improvement
than the architecture comparisons in Section 4.4. This suggests the gap is at
least partly a representation problem that architectural work has been indirectly
120 compensating for.

2.4. Evaluation Methodology in this Area

The methodological pitfalls we control for are documented in exactly the
literature above, which is why we treat them as requirements rather than refine-
ments. Duplication across and within corpora inflates results unless removed
125 before splitting [3, 14]. Project mass is heavily concentrated, so a nominally
cross-project split can degenerate into evaluation on a single repository unless
large projects are handled explicitly. Threshold-dependent metrics are unstable
when the calibration and test distributions differ, which they do by construction
in cross-project evaluation. And functions within a repository are not indepen-
130 dent, so confidence intervals computed over functions understate uncertainty.
Section 4.1.3 states how each is addressed.

2.5. Federated and Privacy-Preserving Variants

A body of work applies federated learning to software engineering tasks,
including defect prediction [19], code clone detection and defect prediction
135 jointly [20], and vulnerability detection [21, 22], typically with token or met-
ric features; FedProto [23] introduced sharing class prototypes rather than pa-
rameters, and differential privacy for learned models follows DP-SGD [24] with
Rényi accounting [25]. Because our abstraction removes project-specific vocab-
ulary from every statistic computed over the graph, it composes naturally with
140 such settings, and we report that as a corollary in Section 5. We do not claim a
federated contribution: in our own experiments ordinary parameter averaging,
without privacy constraints, remained the stronger aggregation strategy.

3. Operation Abstraction

The object of study is a labelling function that replaces a program’s project-specific lexical surface with a small alphabet of operation types fixed by the language grammar. This section specifies it, states the projection lattice relating its three granularities, and defines the graph on which it operates. Section 4 then evaluates it twice: as a static analysis, against annotated ground truth, and as a representation, by its effect on cross-project detection.

3.1. Program Graphs

Each function is parsed with `tree-sitter`, an incremental, error-tolerant parser that recovers a usable tree from non-preprocessed C. That property matters here: vulnerability corpora ship isolated functions without their headers, so a parser requiring a complete translation unit cannot be used. From the resulting tree we build $G = (\mathcal{V}, \mathcal{E})$:

- $\mathcal{V}$: named AST nodes in pre-order, capped at 200;
- $\mathcal{E}_{AST}$: parent-child edges, both directions;
- $\mathcal{E}_{SIB}$: consecutive named siblings, giving source order;
- $\mathcal{E}_{DU}$: a def-use approximation linking successive occurrences of the same identifier leaf.

with $\mathcal{E} = \mathcal{E}_{AST} \cup \mathcal{E}_{SIB} \cup \mathcal{E}_{DU}$, giving on average ~ 130 nodes and ~ 350 edges per function. We state plainly that this is a syntax graph with a data-dependence approximation, not a Code Property Graph [6]: interprocedural control flow and points-to-precise dependence are absent. Section 5.4 discusses what that bounds.

3.2. The Labelling Function ϕ

Let K be the node-kind alphabet of the tree-sitter C grammar ($|K| = 363$) and $\mathcal{T}_{32}$ the operation alphabet of Table 1. The abstraction is a total function

$$\phi : \mathcal{V} \longrightarrow \mathcal{T}_{32} \cup \{\perp\}, \quad (1)$$

where $\perp$ (UNKNOWN) labels nodes carrying no operation semantics. ϕ is deterministic and needs no name resolution, no build system and no whole-program context. It is defined by cases on the node's grammar kind:

1. **Call expressions** are labelled by their callee name through ψ (Section 3.2.1); an unrecognised callee yields `CALL_SITE`. User-defined functions, macros and calls through function pointers therefore all receive a label without resolution.
2. **Kind-determined nodes** follow a fixed table: `subscript_expression` $\rightarrow$ `ARRAY_INDEX`; `field_expression` $\rightarrow$ `FIELD_ACCESS`; `cast_expression` $\rightarrow$ `CAST`; `if/switch/case` $\rightarrow$ `BRANCH_*`; `loops` $\rightarrow$ `LOOP_*`; `break/continue/goto/return` $\rightarrow$ `JUMP_*/RETURN`; `assignment_expression` and initialising declarators $\rightarrow$ `ASSIGN`; `update_expression` $\rightarrow$ `ARITH_ADD`.
3. **Operator-determined nodes**: a `binary_expression` is labelled by its operator token (`+ -` $\rightarrow$ `ARITH_ADD`; `* /` $\rightarrow$ `ARITH_MUL`; `%` $\rightarrow$ `ARITH_MOD`; `bitwise` $\rightarrow$ `BITWISE_OP`; `shifts` $\rightarrow$ `SHIFT_OP`; `equality` $\rightarrow$ `COMPARISON_EQ`; `relational` $\rightarrow$ `COMPARISON_REL`; `logical` $\rightarrow$ `LOGICAL_OP`). A `pointer_expression` is `PTR_DEREF` or `ADDR_OF` by its prefix operator; unary `!` is `LOGICAL_OP`.
4. **Otherwise** $\phi(v) = \perp$.

3.2.1. Callee classification ψ

ψ maps a callee name to an operation class in two stages.

Stage 1 (exact). A fixed allow-list of C standard-library names, grouped by the memory effect they carry: allocation, reallocation, release, block copy, block set, string copy, concatenation, buffer-writing formatting, length query, and stream I/O.

Stage 2 (component). Production C rarely calls those directly: FFmpeg wraps allocation as `av_malloc`, GLib as `g_malloc`, the Linux kernel as `kmalloc`, OpenSSL as `OPENSSL_malloc`. Stage 2 splits the callee identifier into components on underscore and camel-case boundaries and fires when a *whole component* matches an operation word.

Whole-component matching, rather than substring matching, is load-bearing. Substring matching — the obvious implementation — is measurably wrong on
 200 real code: it labels `__put_user` and `skb_put` (writes) as release operations via “put”, `gen_new_label` as allocation via “new”, `float64_is_zero` (a predicate) as a block set via “zero”, and `reallocate_view` as reallocation, while *missing* `copy_to_user`, a genuine bounded copy. Section 4.2 quantifies the resulting error.

205 We further separate stream I/O from buffer formatting. `printf` and `fprintf` write to a stream and carry no caller-buffer overflow semantics; only the `s*printf` and `scanf` families fill a caller-supplied buffer. The distinction is not cosmetic: in our Devign sample `fprintf` and `printf` alone account for 655 call sites, every one of which a conflated rule would count as a memory operation.

210 3.3. The Projection Lattice

Three granularities are defined — $\mathcal{T}_{32}$ (finest), $\mathcal{T}_{16}$, $\mathcal{T}_8$ — related by surjections

$$\pi_{16} : \mathcal{T}_{32} \rightarrow \mathcal{T}_{16}, \quad \pi_8 : \mathcal{T}_{16} \rightarrow \mathcal{T}_8, \quad (2)$$

each extended to fix $\perp$. Coarser labellings are compositions, $\phi_{16} = \pi_{16} \circ \phi$ and $\phi_8 = \pi_8 \circ \pi_{16} \circ \phi$, so the granularities are nested *by construction* rather than
 215 separately specified.

Proposition 1 (Nesting). *For all $u, v \in \mathcal{V}$, $\phi(u) = \phi(v) \Rightarrow \phi_{16}(u) = \phi_{16}(v) \Rightarrow \phi_8(u) = \phi_8(v)$. Equivalently, the partition of $\mathcal{V}$ induced by ϕ_8 is coarser than that induced by ϕ_{16} , which is coarser than that induced by ϕ .*

Proof. Immediate from Eq. (2): ϕ_{16} and ϕ_8 are compositions of ϕ with total
 220 functions, and applying a function to equal arguments yields equal results. The induced partitions are the preimages of a chain of surjections, hence a chain of coarsenings. $\square$

The consequence is that varying granularity varies exactly one quantity — the coarseness of a single labelling — so the granularity experiment in Section 4

Table 1: The $|\mathcal{T}_8| = 8$ operation alphabet and the $\mathcal{T}_{32}$ classes each abstracts. $\perp$ (UNKNOWN) labels nodes with no operation semantics: declarations, identifiers, literals, punctuation.

$\mathcal{T}_8$	abstracts ($\mathcal{T}_{32}$)	carries
MEMORY_ACCESS	MEMORY_{ALLOC,REALLOC,FREE,COPY,SET}, STRING_{COPY,CONCAT,FORMAT,LENGTH}	buffer lifetime, bounds
ARRAY_INDEX	ARRAY_INDEX	indexed access
PTR_DEREF	PTR_DEREF, ADDR_OF, FIELD_ACCESS, CAST	indirection
CONTROL_BRANCH	BRANCH_{IF,SWITCH}, LOOP_{FOR,WHILE}, JUMP_{BREAK,GOTO}, RETURN	reachability
ARITH_OP	ARITH_{ADD,MUL,MOD}, BITWISE_OP, SHIFT_OP	value computation
COMPARISON	COMPARISON_{EQ,REL}, LOGICAL_OP	guards
CALL_SITE	CALL_SITE, IO_CALL	unresolved or non-memory calls
ASSIGN	ASSIGN	value flow

225 is interpretable, which a comparison of three independently designed schemes would not be.

3.4. Node Features

A node is represented as the sum of two embeddings, both over *project-invariant* alphabets:

$$\mathbf{x}_v = E_{\text{op}}[\phi_8(v)] + E_{\text{kind}}[\kappa(v)], \quad (3)$$

230 where $\kappa(v) \in K$ is the grammar kind and $E_{\text{op}}, E_{\text{kind}}$ are learned tables with $|\mathcal{T}_8| + 1 = 9$ and $|K| + 1 = 364$ rows. Both alphabets are fixed by the C grammar and the taxonomy, hence identical across projects and corpora. That invariance is the property we claim explains transfer, and it is what Section 4 tests against a lexical baseline under an otherwise identical protocol.

235 One caveat, stated precisely. Node *features* contain no identifier, but ϕ *consults* identifier text to classify a callee, so a bounded amount of name-derived information — at most $\log_2(|\mathcal{T}_{32}| + 1) \approx 5$ bits per call node — influences the representation. What does not survive is the identifier itself, its spelling, or any project-specific vocabulary.

240 4. Evaluation

We evaluate the abstraction twice, and report negative results alongside positive ones:

- **RQ1 (as an analysis).** How accurately does ϕ label operations, per class, and what does component-based recognition of wrapped APIs buy and cost relative to exact allow-listing?
245
- **RQ2 (as a representation).** With graph, backbone and training budget held fixed, how much does replacing lexical node features with the abstraction change cross-project detection?
- **RQ3 (generality).** Does the effect hold across corpora and across message-passing schemes?
250
- **RQ4 (what does not work).** Do a learned edge-weighting module, an affinity-weighted federated aggregation, or a finer taxonomy improve on their ablations?

4.1. Setup

255 4.1.1. Corpora

We use BigVul [26], DiverseVul [14] and PrimeVul [3], loaded from pinned public mirrors. Each is sampled to 8,000 functions. We disclose two properties that are usually left implicit and that materially affect cross-project claims: the sample is taken in corpus order rather than at random, and project mass is highly concentrated — in BigVul, Chrome and `linux` together supply about
260 two thirds of the functions.

4.1.2. Deduplication

Near-duplicate functions are removed *before* any split. Vended code appears under several project names, so without this step a held-out function can
265 have a twin in training and project-disjointness is defeated silently. Signatures are computed over the abstracted node sequence, which catches renamed copies that text hashing misses; 2–3% of functions are removed.

4.1.3. Cross-project protocol

Evaluation is repeated project-level GroupKFold: projects are dealt into
270 five folds, no project is split, and each fold is used once as the test set. Two

provisions matter. First, projects too large to be placed without dominating a fold — Chrome and `linux` — are held *train-only* and reported as such, so that no fold labelled “cross-project” is in fact an evaluation on one repository. Second, all vocabularies and label bucketings are fitted on training projects only. The
275 resulting folds are stable: test fraction 0.071 on every fold, 39–49 test projects per fold.

We report AUC as the primary metric and AUPRC as the imbalanced-data companion, both threshold-free. F1 is reported only against an explicit floor: the trivial classifier that labels every function vulnerable achieves $F_1 = 0.161$ at
280 the 8.8% prevalence of these test sets, which is above several F1 values reported as results in the literature.

4.1.4. Statistical treatment

Fold-level means are reported with standard deviations, but inference uses a *cluster bootstrap* that resamples held-out projects with replacement (1000 re-
285 samples), because functions within a repository share style, APIs and often near-duplicate code and are therefore not independent. We report the paired difference between systems with a 95% percentile interval and a two-sided bootstrap p -value. Where an interval crosses zero we say so and draw no conclusion.

4.2. RQ1: ϕ as a static analysis

290 We extract every callee name in the Devign corpus, sample 212 names stratified over predicted classes (a precision arm) together with a frequency-weighted sample of unmatched names (a recall arm), and annotate the true class from documented API semantics. The annotation criteria are stated in the released script so that a reader can contest a specific label rather than an opaque score:
295 an allocation label requires that the callee’s documented purpose is to obtain, resize or release a heap allocation the *caller* then owns; a copy or set label requires that it writes a caller-supplied destination of caller-specified length; `printf`-family calls that write to a stream are I/O, not buffer formatting.

Table 2: RQ1: ϕ evaluated as a static analysis on 212 annotated callee names from Devign. Call-site-weighted accuracy is 0.976; macro precision 0.976, macro recall 0.938.

Class	Precision	Recall	F_1	tp/fp/fn
MEMORY_ALLOC	0.950	0.905	0.927	19/1/2
MEMORY_REALLOC	1.000	1.000	1.000	10/0/0
MEMORY_FREE	0.950	1.000	0.974	19/1/0
MEMORY_COPY	0.857	0.600	0.706	6/1/4
MEMORY_SET	1.000	1.000	1.000	2/0/0
STRING_COPY	1.000	1.000	1.000	5/0/0
STRING_CONCAT	1.000	1.000	1.000	4/0/0
STRING_FORMAT	1.000	1.000	1.000	7/0/0
STRING_LENGTH	1.000	1.000	1.000	2/0/0
IO_CALL	1.000	0.880	0.936	22/0/3
Macro	0.976	0.938	—	

4.2.1. What the validation changed

300 The exercise was not confirmatory. It changed the analysis three times, and each change is a result about how such abstractions should be built.

Substring matching is unsafe. Recognising wrapped allocators by substring — the obvious implementation — labels `__put_user` and `skb_put`, both *writes*, as release operations via “put”; `gen_new_label` as an allocation via “new”; 305 `float64_is_zero`, a predicate, as a block set via “zero”; and `reallocate_view` as a reallocation — while missing `copy_to_user`, a genuine bounded copy. Matching whole identifier components removes all of these.

Stream I/O is not buffer formatting. Placing `printf` and `fprintf` in the formatting class causes 655 call sites in this corpus alone to be counted 310 as memory operations. Only the `sprintf` and `scanf` families write into a caller-supplied buffer.

A precision/recall trade-off, made explicit. Table 3 reports the copy rule at two operating points. A permissive needle matching any “copy” com-

Table 3: RQ1: the copy rule at two operating points, showing the precision/recall trade-off the gold set exposed.

Rule	MEMORY_COPY			Macro-P	Weighted acc.
	P	R	F_1		
Permissive (any “copy”)	0.450	0.900	0.600	0.935	0.971
Precise (adopted)	0.857	0.600	0.706	0.976	0.976

ponent reaches recall 0.900 but precision 0.450, because it captures SIMD lane
 315 extraction (`__msa_copy_u_w`) and pixel-format macros (`COPY16T09_OR_10`), nei-
 ther of which writes a caller buffer. Requiring an explicit copy primitive, or
 “copy” adjacent to a direction word, raises precision to 0.857 at the cost of re-
 call 0.600 — it now misses `AV_COPY32` and `av_dict_copy`. We adopt the precise
 variant, because a false memory-operation label injects a spurious safety-relevant
 320 node into the graph whereas a miss merely leaves the node as a generic call site.

4.2.2. Coverage and residual disagreements

Component matching raises coverage of call sites from 4.7% under exact
 allow-listing to 10.6%. Ten disagreements remain in the sample, and we name
 the informative ones rather than aggregating them away. `bdrv_pread` (36
 325 sites) is a positioned block-device read that no lexical rule recognises as I/O.
`g_strdup_printf` both allocates and formats, which a single-label taxonomy
 cannot express; we label it as allocation, the stronger memory effect. `tcg_temp_free`
 releases an intermediate-representation temporary in QEMU’s own allocator dis-
 cipline rather than heap memory — a release, but not the kind the taxonomy
 330 was designed around. These are limitations of a name-based analysis without
 type information, and they bound what the abstraction can claim.

4.3. RQ2: the representation effect

Everything except the node features is held fixed: graph topology, backbone
 (GGNN), depth, hidden width, mean-max readout, classification head, opti-

Table 4: RQ2: node-feature mode under an otherwise identical protocol (BigVul, 10 project-level folds, mean $\pm$ std). Differences are paired per fold; intervals are cluster bootstrap over held-out projects.

Node features	AUC	AUPRC
Trivial (all-positive)	0.500	0.088
Lexical token	0.587 ± 0.050	0.125 ± 0.040
Operation + grammar kind	0.668 ± 0.027	0.164 ± 0.034
Paired difference	+0.064	+0.022
95% CI (cluster bootstrap)	[+0.028, +0.112]	[+0.000, +0.049]
p (two-sided)	0.001	0.047
Folds won	8/10	8/10

335 miser, class weighting and gradient budget. Table 4 reports BigVul over ten folds.

The abstraction improves cross-project ranking by +0.064 AUC ($p = 0.001$) and +0.022 AUPRC ($p = 0.047$), winning 8 of 10 folds on both metrics, with both intervals excluding zero. It is also markedly more stable across folds (AUC
340 standard deviation 0.027 versus 0.050), which is what one expects if the lexical model’s fold-to-fold variation is driven by how much of the test vocabulary it happened to see in training.

Two honest qualifications. Absolute performance remains modest: AUPRC of 0.164 against a base rate of 0.088 is a lift of roughly $1.9\times$, and cross-project
345 detection is not solved by this or, on the present evidence, by any published method. And the AUPRC interval’s lower bound is +0.000: the effect on precision at the top of the ranking is real but small, and we do not claim more.

4.4. RQ3: generality

The same contrast is run on a second corpus and with two further message-
350 passing schemes (GAT, GIN) to test whether the effect is specific to one dataset or one architecture. Results are reported in Table ??; the design is identical to

RQ2 in every respect other than the factor named.

4.5. RQ4: what does not work

We report three negative results at the same length as the positive one,
355 because each corresponds to a mechanism we implemented and expected to
help.

Learned edge weighting. A module that learns a soft importance mask
over edges, trained jointly with the classifier and regularised toward sparsity,
does not improve on its own ablation; removing it leaves AUC unchanged and
360 slightly improves F1.

Affinity-weighted aggregation. In a federated variant, weighting client
contributions by the cosine similarity of their class prototypes does not improve
on uniform averaging; the ablation without it wins on the majority of seeds.

Taxonomy refinement. Increasing the operation alphabet from 8 to 16
365 or 32 types (Proposition 1 guarantees these are strict refinements of the same
labelling) does not produce a corresponding gain. The likely reason is visible in
the class distribution: the memory-operation class that motivates the taxonomy
fires on under 1% of AST nodes, so subdividing it cannot move an aggregate
metric.

370 5. Discussion

5.1. Why the abstraction transfers

The result is consistent with a simple account. A lexical model’s features
are drawn from the vocabulary of the projects it trained on; at test time on
an unseen repository, much of its input is out-of-vocabulary, and what remains
375 correlates with project identity rather than with the weakness. The abstracted
model’s alphabet is fixed by the C grammar and the operation taxonomy, so its
input distribution is, by construction, the same on every project. The stability
we observe supports this: fold-to-fold AUC deviation is roughly half that of
the lexical model, which is what one expects if lexical variance is driven by
380 accidental vocabulary overlap between train and test.

This also explains why the effect is larger on ranking (AUC) than at the top of the ranking (AUPRC). The abstraction supplies a consistent notion of *what kind of operation* a node performs, which is enough to order functions by risk; it discards the specific API, which is often what distinguishes a genuinely
385 dangerous call from a benign one of the same type.

5.2. *Validating an abstraction is not optional*

The most transferable methodological point concerns RQ1. Program abstractions are routinely introduced as preprocessing and never evaluated: a taxonomy is listed, a mapping asserted, and downstream accuracy is reported.
390 Our experience argues this is unsafe. Two defects in our own mapping — substring matching that classified writes and predicates as memory operations, and a formatter class that absorbed 655 stream-I/O call sites — were invisible to every downstream metric and were caught only by annotating call sites and computing per-class precision. Both would have inflated the very statistic we
395 use to motivate the taxonomy.

The cost is modest: a few hundred annotated names. We would encourage the practice as a default for work that introduces a code abstraction, in the same way that a static analysis is expected to report precision and recall.

5.3. *Negative results*

400 We implemented a learned edge-weighting module and an affinity-weighted federated aggregation, and expected both to help. Neither improves on its ablation. We report them because the alternative — keeping components a paper names but its data does not support — is how unreproducible results enter the literature, and because the negative result is itself informative: it locates
405 the effect in the representation rather than in the machinery built around it.

The taxonomy-refinement result has a concrete explanation worth recording. The memory-operation class fires on under 1% of AST nodes, so subdividing it into allocation, release, copy and set cannot shift an aggregate metric however well-motivated the distinction is semantically. A finer taxonomy would need to
410 be paired with an evaluation that weights the nodes it refines.

5.4. Limitations

Graph fidelity. Our graphs are tree-sitter syntax trees with sibling and def-use approximation edges, not Code Property Graphs: interprocedural control flow and points-to-precise dependence are absent. Every system we compare
415 consumes the same graphs, so the comparison is internally fair, but the absolute ceiling is bounded by this choice.

Scale and language. Experiments use 8,000 functions per corpus, compact models, and C only. The abstraction is specified over a tree-sitter grammar and is therefore instantiable for other languages, but we have not demonstrated that,
420 and a taxonomy validated on C should not be assumed to transfer to, say, Rust’s ownership discipline.

Sampling and concentration. Corpus samples are taken in file order rather than at random, and project mass is concentrated (Chrome and `linux` supply about two thirds of BigVul). We hold those projects train-only and
425 report it, but a differently constructed sample could yield different absolute numbers.

Label noise. We inherit the labels of the underlying corpora, whose accuracy is itself contested in the literature that produced them.

Effect size. The AUPRC improvement is small and its interval’s lower
430 bound is at zero. The claim we defend is that the representation matters more than the architecture choices it is usually compared against — not that cross-project detection is solved.

5.5. A corollary for federated deployment

Because statistics computed over the abstracted graphs contain no project-specific vocabulary, class-conditioned summaries can be shared between organi-
435 sations without transmitting code, and perturbed to satisfy differential privacy at low dimensionality. We report this as a corollary rather than a contribution: in our experiments, ordinary parameter averaging without any privacy constraint remained the stronger aggregation strategy, and we found no ev-

440 idence that the prototype-based alternative beats it absent a formal privacy requirement.

6. Conclusion

We asked how much of the cross-project gap in vulnerability detection is attributable to the program representation rather than the model, and answered
445 it by specifying an abstraction and measuring it under controls.

ϕ is a total labelling from a C grammar’s node alphabet onto an operation taxonomy, with a projection lattice that makes its granularities nested by construction. Evaluated as a static analysis it exposes failure modes that downstream accuracy hides; evaluated as a representation, with graph, backbone,
450 depth, readout and training budget all held fixed, it improves cross-project ranking by +0.064 AUC ($p = 0.001$) and +0.022 AUPRC ($p = 0.047$), winning 8 of 10 project-level folds, and roughly halves fold-to-fold variance.

We are equally clear about what we did not find. A learned edge-weighting module, an affinity-weighted federated aggregation, and taxonomy refinement
455 beyond eight types all fail to improve on their ablations. Absolute performance stays modest — an AUPRC lift of about $1.9\times$ over base rate — and cross-project detection remains open.

The methodological suggestion we would most like to see adopted is the cheapest one: when a paper introduces a code abstraction, evaluate it as an
460 analysis. A few hundred annotated call sites caught two defects in our own mapping that no downstream metric revealed, and that would otherwise have inflated the statistic used to motivate the design.

References

- [1] S. Chakraborty, R. Krishna, Y. Ding, B. Ray, Deep learning based vulnerability detection: Are we there yet?, IEEE Transactions on Software
465 Engineering 48 (9) (2022) 3280–3296. doi:10.1109/TSE.2021.3087402.

- [2] P. Chakraborty, K. K. Arumugam, M. Alfadel, M. Nagappan, S. McIntosh, Revisiting the performance of deep learning-based vulnerability detection on realistic datasets, *IEEE Transactions on Software Engineering* 50 (8) (2024) 2163–2177, arXiv:2407.03093. doi:10.1109/TSE.2024.3423712.
- [3] Y. Ding, Y. Fu, O. Ibrahim, C. Sitawarin, X. Chen, B. Alomair, D. Wagner, B. Ray, Y. Chen, Vulnerability detection with code language models: How far are we?, in: *Proceedings of the 47th IEEE/ACM International Conference on Software Engineering (ICSE)*, IEEE, 2025, pp. 469–481, arXiv:2403.18624. doi:10.1109/ICSE55347.2025.00038.
- [4] Z. Li, D. Zou, S. Xu, X. Ou, H. Jin, S. Wang, Z. Deng, Y. Zhong, VulDeePecker: A deep learning-based system for vulnerability detection, in: *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, Internet Society, 2018. doi:10.14722/ndss.2018.23158.
- [5] Z. Li, D. Zou, S. Xu, H. Jin, Y. Zhu, Z. Chen, SySeVR: A framework for using deep learning to detect software vulnerabilities, *IEEE Transactions on Dependable and Secure Computing* 19 (4) (2022) 2244–2258. doi:10.1109/TDSC.2021.3051427.
- [6] F. Yamaguchi, N. Golde, D. Arp, K. Rieck, Modeling and discovering vulnerabilities with code property graphs, in: *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*, IEEE, 2014, pp. 590–604. doi:10.1109/SP.2014.44.
- [7] Y. Zhou, S. Liu, J. K. Siow, X. Du, Y. Liu, Devign: Effective vulnerability identification by learning comprehensive program semantics via graph neural networks, in: *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32, 2019, arXiv:1909.03496.
- [8] X. Cheng, H. Wang, J. Hua, G. Xu, Y. Sui, DeepWukong: Statically detecting software vulnerabilities using deep graph neural networks, *ACM Transactions on Software Engineering and Methodology (TOSEM)* 30 (3) (2021) 1–33. doi:10.1145/3436877.

- [9] Y. Li, S. Wang, T. N. Nguyen, Vulnerability detection with fine-grained interpretations, in: Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE), ACM, 2021, pp. 292–303. doi:10.1145/3468264.3468597.
- [10] Y. Luo, W. Xu, D. Xu, Detecting code vulnerabilities with heterogeneous GNN training, International Journal of Information Security 24, arXiv:2502.16835 (2025). doi:10.1007/s10207-025-01132-x.
- [11] M. Fu, C. Tantithamthavorn, LineVul: A transformer-based line-level vulnerability prediction, in: Proceedings of the 19th International Conference on Mining Software Repositories (MSR), ACM, 2022, pp. 608–620. doi:10.1145/3524842.3527927.
- [12] D. Hin, A. Kan, H. Chen, M. A. Babar, LineVD: Statement-level vulnerability detection using graph neural networks, in: Proceedings of the 19th International Conference on Mining Software Repositories (MSR), ACM, 2022, pp. 596–607. doi:10.1145/3524842.3527949.
- [13] R. Liu, Y. Wang, H. Xu, J. Sun, F. Zhang, P. Li, Z. Guo, VulLMGNNs: Fusing language models and online-distilled graph neural networks for code vulnerability detection, Information Fusion 115 (2025) 102748, arXiv:2404.14719. doi:10.1016/j.inffus.2024.102748.
- [14] Y. Chen, Z. Ding, L. Alowain, X. Chen, D. Wagner, DiverseVul: A new vulnerable source code dataset for deep learning based vulnerability detection, in: Proceedings of the 26th International Symposium on Research in Attacks, Intrusions and Defenses (RAID), ACM, 2023, pp. 654–668. doi:10.1145/3607199.3607242.
- [15] S. Liu, G. Lin, L. Qu, J. Zhang, O. D. Vel, P. Montague, Y. Xiang, CD-VulD: Cross-domain vulnerability discovery based on deep domain adaptation, IEEE Transactions on Dependable and Secure Computing 19 (1) (2022) 438–451. doi:10.1109/TDSC.2020.2984505.

- 525 [16] C. Zhang, B. Liu, Y. Xin, L. Yao, CPVD: Cross project vulnerability detection based on graph attention network and domain adaptation, *IEEE Transactions on Software Engineering* 49 (8) (2023) 4152–4168. doi:10.1109/TSE.2023.3285910.
- [17] X. Li, Y. Xin, H. Zhu, Y. Yang, Y. Chen, Cross-domain vulnerability detection using graph embedding and domain adaptation, *Computers & Security* 530 125 (2023) 103017. doi:10.1016/j.cose.2022.103017.
- [18] Z. Cai, Y. Cai, X. Chen, G. Lu, W. Pei, J. Zhao, CSVD-TF: Cross-project software vulnerability detection with TrAdaBoost by fusing expert metrics and semantic metrics, *Journal of Systems and Software* 213 (2024) 112038. doi:10.1016/j.jss.2024.112038. 535
- [19] H. Yamamoto, D. Wang, G. K. Rajbahadur, M. Kondo, Y. Kamei, N. Ubayashi, Towards privacy preserving cross project defect prediction with federated learning, in: *Proceedings of the 30th IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, IEEE, 2023, pp. 485–496. doi:10.1109/SANER56733.2023.00052. 540
- [20] Y. Yang, X. Hu, Z. Gao, J. Chen, C. Ni, X. Xia, D. Lo, Federated learning for software engineering: A case study of code clone detection and defect prediction, *IEEE Transactions on Software Engineering* 50 (2) (2024) 296–321. doi:10.1109/TSE.2023.3347898.
- 545 [21] C. Zhang, T. Yu, B. Liu, Y. Xin, Vulnerability detection based on federated learning, *Information and Software Technology* 167 (2024) 107371. doi:10.1016/j.infsof.2023.107371.
- [22] P. Zhou, M. Hu, X. Xie, Y. Liu, M. Chen, VulFL: An empirical study of vulnerability detection using federated learning, *arXiv preprint-ArXiv:2411.16099* (2024). 550
- [23] Y. Tan, G. Long, L. Liu, T. Zhou, Q. Lu, J. Jiang, C. Zhang, FedProto: Federated prototype learning across heterogeneous clients, in: *Proceedings*

of the 36th AAAI Conference on Artificial Intelligence, 2022, pp. 8432–8440.
doi:10.1609/aaai.v36i8.20819.

- 555 [24] M. Abadi, A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar,
L. Zhang, Deep learning with differential privacy, in: Proceedings of the
2016 ACM SIGSAC Conference on Computer and Communications Secu-
rity (CCS), ACM, 2016, pp. 308–318. doi:10.1145/2976749.2978318.
- [25] I. Mironov, K. Talwar, L. Zhang, Rényi differential privacy of the sampled
560 gaussian mechanism, in: arXiv preprint, 2019, arXiv:1908.10530.
- [26] J. Fan, Y. Li, S. Wang, T. N. Nguyen, A C/C++ code vulnerability dataset
with code changes and CVE summaries, in: Proceedings of the 17th Inter-
national Conference on Mining Software Repositories (MSR), ACM, 2020,
pp. 508–512. doi:10.1145/3379597.3387501.